


Constraint duals and KL coefficient (5-seed mean $\pm$ std $\cdot$ training logs)

(a)

Adaptive Lagrange multipliers



(b)

KL regularization schedule (not KL divergence)

